



Choose scale: small, medium, large or xlarge

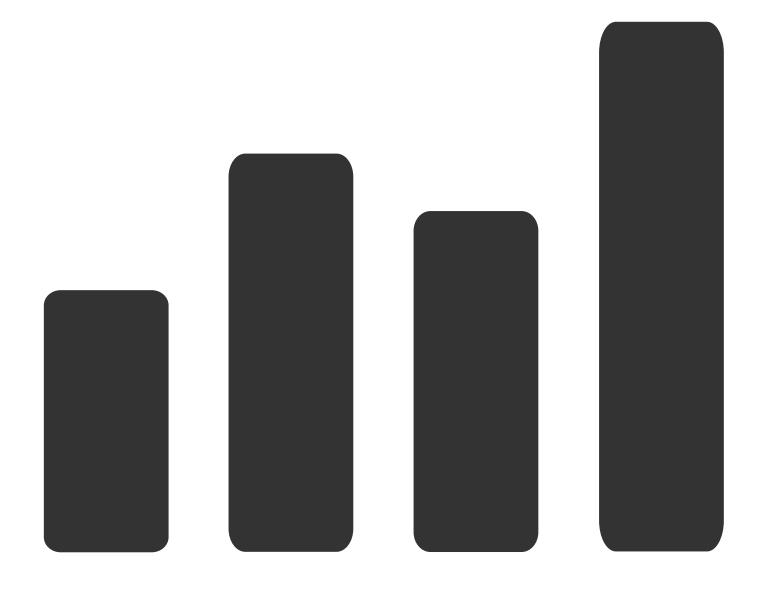






Train a CLIP model with a fixed architecture and hyper-parameters





Evaluate the model on 38 zero-shot downstream tasks